

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/222037642>

A comprehensive review and evaluation of permutation flowshop heuristics to minimize flowtime

Article in *European Journal of Operational Research* · September 2005

DOI: 10.1016/j.ejor.2004.04.017

CITATIONS

557

READS

1,109

2 authors:



[Rubén Ruiz](#)

Universitat Politècnica de València

192 PUBLICATIONS 12,109 CITATIONS

SEE PROFILE



[Concepción Maroto](#)

Universitat Politècnica de València

60 PUBLICATIONS 3,269 CITATIONS

SEE PROFILE

A Comprehensive Review and Evaluation of Permutation Flowshop Heuristics

Rubén Ruiz*, Concepción Maroto

Dpto. de Estadística e Investigación Operativa Aplicadas y Calidad

Universidad Politécnica de Valencia

6th January 2004

Abstract

In this work we present a review and comparative evaluation of heuristics and metaheuristics for the well known permutation flowshop problem with the makespan criterion. A number of reviews and evaluations have already been proposed. However, the evaluations do not include the latest heuristics available and there is still no comparison of metaheuristics. Furthermore, since no common benchmarks and computing platforms are used, the results cannot be generalised. We propose a comparison of 25 methods, ranging from the classical Johnson's algorithm or dispatching rules to the most recent metaheuristics, including tabu search, simulated annealing, genetic algorithms, iterated local search and hybrid techniques. For the evaluation we use the

*Corresponding author: Rubén Ruiz García. Universidad Politécnica de Valencia. Departamento de Estadística e Investigación Operativa Aplicadas y Calidad. Camino de Vera S/N, 46021, Valencia, SPAIN. Tel: +34 96 387 70 07, ext: 74946. Fax: +34 96 387 74 99. e-mail: rruiz@eio.upv.es

standard test of Taillard (1993) composed of 120 instances of different sizes. In the evaluations we use the experimental design approach to obtain valid conclusions on the effectiveness and efficiency of the different methods tested.

Keywords: Flowshop, Scheduling, Heuristics, Metaheuristics.

1 Introduction

In a flowshop scheduling problem there is a set of n jobs, tasks or items $(1, \dots, n)$ to be processed in a set of m machines or processors $(1, \dots, m)$ in the same order, i.e. first in machine 1 then on machine 2 and so on until machine m . The objective is to find a sequence for the processing of the jobs in the machines so that a given criterion is optimised. In the literature, the most common criterion is the minimisation of the total completion time or makespan of the schedule (C_{max}), this is sometimes referred to as maximum flow time or F_{max} . The processing times needed for the jobs on the machines are noted as p_{ij} , where $i = 1, \dots, n$ and $j = 1, \dots, m$, these times are fixed, known in advance and nonnegative. There are several assumptions that are commonly made regarding this problem (Baker, 1974):

- Each job i can be processed at most on one machine j at the same time.
- Each machine m can process only one job i at a time.
- No preemption is allowed, i.e. the processing of a job i on a machine j can not be interrupted.
- All jobs are independent and are available for processing at time 0.

- The set-up times of the jobs on machines are negligible and therefore can be ignored.
- The machines are continuously available.
- In-process inventory is allowed. If the next machine on the sequence needed by a job is not available, the job can wait and joins the queue at that machine.

This problem was initially classified as $n/m/F/F_{max}$ following the four parameter notation $A/B/C/D$ by Conway et al. (1967) and then as $F//C_{max}$ following the three parameter notation $\alpha/\beta/\gamma$ introduced by Graham et al. (1979). Although the flowshop scheduling problem is solvable to optimality in polynomial time when $m = 2$ (Johnson, 1954), it is known to be \mathcal{NP} -Complete in the strong sense when $m \geq 3$ (see Garey et al., 1976). In general $(n!)^m$ schedules have to be considered. This is why the problem is further restricted in the literature by not allowing *job passing*, i.e. the processing sequence of the jobs is the same for all the machines, so no job can pass another in the following machine. In this case “only” $n!$ schedules have to be considered and the problem is then known as *permutation flowshop* and is classified as $n/m/P/F_{max}$ or as $F/prmu/C_{max}$ (see Pinedo, 2002). In this work we will concentrate on this last type of flowshop environment.

In the literature of flowshop scheduling there are various reviews (for example, Dudek et al., 1992 or Reisman et al., 1994), and also there are some comparative evaluations available. In Ponnambalam et al. (2001), five different heuristics are compared against only 21 typical test problems. Turner and Booth (1987) compare two famous heuristics with a set of 350 random problems. In the work of Byung Park et al. (1984), a total of 16 heuristics are compared against 1,500 random generated problems of a maximum size of 30 jobs and 20 machines. Dannenbring (1977) evaluated 11 different

heuristics with a benchmark set of 1,580 random problems of various sizes and characteristics. Finally, Ashour (1970), compared some of the earliest methods for the permutation flowshop problem.

As we can see, all these evaluations have several shortcomings; usually no common data sets and/or computer mainframes are used for the results to be generalised. The comparisons are usually performed between no more than a few heuristics and, furthermore, no comparison deals with the most recent heuristics available. It is also difficult to find comparative evaluations between advanced metaheuristics such as Simulated Annealing or Genetic Algorithms.

For all aforementioned reasons, today it is impossible to make a global comparative evaluation of the different methods proposed in the literature, since the evaluations are partial and the benchmarks used are not standard, therefore the results are not reproducible.

The objective of this paper is to give an up to date comprehensive review and evaluation of many existing heuristics and metaheuristics for the permutation flowshop. We evaluate a total of 25 heuristics, from the classical methods to the most recent heuristics and advanced metaheuristics. The comparisons are made against Taillard (1993)'s famous benchmarks.

The rest of the paper is organised as follows. In section 2 the most well-known heuristics for the permutation flowshop (PFSP from now on) are reviewed. Section 3 deals with metaheuristics for the PFSP, from Simulated Annealing to complex hybrid techniques. A comparison of the various heuristics and metaheuristics can be seen in Section 4. Finally, in Section 5 some conclusions are given.

2 Heuristics for the flowshop scheduling problem

The complexity of the flowshop scheduling problem renders exact solution methods impractical for instances of more than a few jobs and/or machines. This is the main reason for the various heuristic methods proposed in the literature, some of which will be explained here. The heuristics can be divided in either *constructive heuristics* or *improvement heuristics*, the former are heuristics that build a feasible schedule from scratch and the latter are heuristics that try to improve a previously generated schedule by normally applying some form of specific problem knowledge.

2.1 Constructive Heuristics

Johnson’s algorithm (1954) is the earliest known heuristic for the PFSP, which provides an optimal solution for two machines. Moreover, it can be used as a heuristic for the m machine case by clustering the m machines into two “virtual” machines. The computational complexity of this heuristic is $\mathcal{O}(n \log n)$. Other authors have used the general ideas of Johnson’s rule in their algorithms, for example, Dudek and Teuton (1964) developed an m -stage rule for the PFSP that minimises the idle time accumulated on the last machine when processing each job by using Johnson’s approach. Campbell et al. (1970) developed a heuristic algorithm which is basically an extension of Johnson’s algorithm. In this case, several schedules are constructed and the best one is given as a result. The heuristic is known as CDS and builds $m - 1$ schedules by clustering the m original machines into two virtual machines and solving the generated two machine problem by repeatedly using Johnson’s rule. The CDS heuristic has a computational complexity of $\mathcal{O}(m^2n + mn \log n)$.

In a more recent work, Koulamas (1998) reported a new two phase heuristic, called HFC. In the first phase, the HFC heuristic makes extensive use of Johnson’s algorithm. The second phase improves the resulting schedule from the first phase by allowing job passing between machines, i.e. by allowing non-permutation schedules. This is a very interesting idea, since it is known that permutation schedules are only dominant for the three-machine case. In the general m machine case, a permutation schedule is not necessarily optimal anymore (Potts et al., 1991). Although this heuristic departs from the PFSP problem by allowing job passing, we have included it in the review for comparison reasons since we believe that some instances might benefit from job passing. Taking into account both phases, the general computational complexity of this heuristic is roughly $\mathcal{O}(m^2n^2)$.

Another approach is to assign a weight or “index” to every job and then arrange the sequence by sorting the jobs according to the assigned index. This idea was first exploited by Palmer (1965) when he developed a very simple heuristic in which for every job a “slope index” is calculated and then the jobs are scheduled by non-increasing order of this index, which leads to a computational complexity of $\mathcal{O}(nm + n \log n)$. This idea has been used in later papers, for example Gupta (1971) proposed a modification of Palmer’s slope index which exploited some similarities between scheduling and sorting problems. Similarly, Bonney and Gundry (1976) worked on the idea of using the geometrical properties of the cumulative process times of the jobs and a slope matching method for scheduling the PFSP. Hundal and Rajgopal (1988) proposed a very simple extension to Palmer’s heuristic by exploiting the fact that when m is an odd number, Palmer’s slope index returns the value 0 for the machine $(m + 1)/2$ and consequently ignores it. In order to overcome this, two more slope indexes are calculated and with these two slope indexes and the original Palmer’s slope index, three schedules are cal-

culated and the best one is given as a result. Since the procedure consists in calculating three indexes, the computational complexity is the same of Palmer's $\mathcal{O}(nm + n \log n)$.

Dannenbring (1977)'s Rapid Access (RA) heuristic is a mixture of the previous ideas of Johnson's algorithm and Palmer's slope index. In this case a virtual two machine problem is defined as it is in the CDS heuristic, but instead of directly applying Johnson's algorithm over the processing times, two weighting schemes are calculated, one for each machine, and then Johnson's algorithm is applied. The weighting schemes give the processing times for the jobs in the two virtual machines. As the name of the heuristic implies, RA provides a good solution in a very quick and easy way $\mathcal{O}(mn + n \log n)$. Since all the jobs in a PFSP form a permutation, many proposed methods work with the idea of exchanging the position of the jobs in the sequence or inserting jobs at different locations to obtain better results. Page (1961) proposed three heuristic methods based on the similarities of the PFSP with sorting methods. The idea is to obtain a good ordering of the jobs and to subsequently improve this order by means of job exchanges.

Nawaz et al. (1983)'s NEH heuristic is regarded as the best heuristic for the PFSP (see Taillard, 1990). It is based on the idea that jobs with high processing times on all the machines should be scheduled as early in the sequence as possible. The procedure is straightforward:

1. The total processing times for the jobs are calculated using the formulae: $\forall \text{ job } i, i = 1, \dots, n, \quad P_i = \sum_{j=1}^m p_{ij}$.
2. The jobs are sorted in non-increasing order of P_i . Then the first two jobs (those two with higher P_i) are taken and the two possible schedules containing them are evaluated.
3. Take job $i, i = 3, \dots, n$ and find the best schedule by placing it in all the

possible i positions in the sequence of jobs that are already scheduled. For example, if $i = 4$, the already constructed sequence would contain the first three jobs of the sorted list calculated in step 2, then the fourth job could be placed either in the first, in the second, in the third or in the last position of the sequence. The best sequence of the four would be selected for the next iteration.

As we have seen, the NEH heuristic is based neither on Johnson's algorithm nor on slope indexes. The only drawback is that a total of $[n(n + 1)/2] - 1$ schedules have to be evaluated, being n of those schedules complete sequences. This makes the complexity of NEH rise to $\mathcal{O}(n^3m)$ which can be lengthy for big problem instances. However, Taillard (1990) reduced NEH's complexity to $\mathcal{O}(n^2m)$ by calculating all the partial schedules in a given iteration in a single step.

Sarin and Lefoka (1993) exploited the idea of minimising idle time on the last machine since any increase in the idle time on the last machine will translate into an increase in the total completion time or makespan. In this way, the sequence is completed by inserting one job at a time and priority is given to the job that, once added to the sequence, would result in minimal added idle time on machine m . The method proposed compares well with the NEH heuristic but only when the number of machines in a problem exceeds the number of jobs. Davoud Pour (2001) proposed another insertion method. This new heuristic is based on the idea of job exchanging and is similar to the NEH method. The performance of this method is evaluated against the NEH, CDS and Palmer's heuristics showing better effectiveness only when a big number of machines is considered, and being the computational complexity $\mathcal{O}(n^3m)$. More recently, Framinan et al. (2003) have published a study about the NEH heuristic where different initialisations and orderings are considered. The study also includes different objective functions.

Other authors have proposed heuristics that use one or more of the previous ideas, for example, Gupta (1972) proposed three heuristic methods, named *minimum idle time* (MINIT), *minimum completion time* (MICOT) and MINIMAX algorithms, the first two are based on job pair exchanges and the MINIMAX is based on Johnson's rule. These three algorithms were tested with the objectives of C_{max} and mean flowtime (\bar{F}) and compared with the CDS algorithm, proving to be superior only when considering the \bar{F} objective.

Additionally, there are many other methods available that are not based on Johnson's or Palmer's ideas and that do not construct sequences by job exchanges and/or insertions only. For example, King and Spachis (1980) evaluated various heuristics for the PFSP and for the flowshop with no job waiting (*no-wait flowshop*). For the PFSP, a total of five heuristics based on dispatching rules were developed. A different approach is shown in Stinson and Smith (1982) where the authors solve the permutation flowshop problem by using a well known heuristic for the Travelling Salesman Problem (TSP).

2.2 Improvement Heuristics

Contrary to constructive heuristics, improvement heuristics start from an already built schedule and try to improve it by some given procedure. Dannenbring (1977) proposed two simple improvement heuristics, these are *Rapid Access with Close Order Search* (RACS) and *Rapid Access with Extensive Search* (RAES). The reason behind these two heuristics is that Dannenbring found that simply swapping two adjacent jobs in a sequence obtained by the RA heuristic resulted in an optimal schedule. RACS works by swapping every adjacent pair of jobs in a sequence (this is $n - 1$ steps). The best schedule among the $n - 1$ generated is then given as a result. In RAES heuristic, RACS is repeatedly applied while improvements are found. Both RACS and

RAES heuristics start from a schedule generated with the RA constructive heuristic.

Ho and Chang (1991) developed a method that works with the idea of minimising the elapsed times between the end of the processing of a job in a machine and the beginning of the processing of the same job in the following machine in the sequence. The authors refer to this time as “gap”. The algorithm calculates the gaps for every possible pair of jobs and machines and then by a series of calculations, the heuristic swaps jobs depending on the value of the gaps associated with them. The heuristic starts from the CDS heuristic by Campbell et al.

Recently, Suliman (2000) developed an improvement heuristic. In the first phase, a schedule is generated with the CDS heuristic method. In the second phase, the schedule generated is improved with a job pair exchange mechanism. In order to reduce the computational burden of an exhaustive pair exchange mechanism, a directionality constraint is imposed to reduce the search space. For example, if by moving a job forward, a better schedule is obtained, it is assumed that better schedules can be achieved by maintaining the forward movement and not allowing a backward movement.

3 Metaheuristics for the flowshop scheduling problem

Metaheuristics are general heuristic procedures that can be applied to many problems, and, in our case, to the PFSP. These methods normally start from a sequence constructed by heuristics and iterate until a stopping criterion is met. There is plenty of research work done for the PFSP and metaheuristics. In this section, we will point out a few noteworthy papers mainly dealing with Simulated Annealing (SA), Tabu Search (TS), Genetic Algorithms (GA) and

other metaheuristics, as well as hybrid methods.

Osman and Potts (1989) proposed a simple Simulated Annealing algorithm using a shift neighbourhood and a random neighbourhood search. In the same year Widmer and Hertz (1989) presented a method called SPIRIT, which is a two phase heuristic. In the first phase, a problem is generated with an analogy with the Open Travelling Salesman Problem (OTSP) and then, this problem is solved with an insertion method to obtain an initial solution. In the second phase, a Tabu Search metaheuristic with standard parameters and exchange neighbourhood is used to improve the incumbent solution. Taillard (1990) also presented a similar procedure to that of Widmer and Hertz. A Tabu Search technique is applied to a schedule generated by an improved NEH heuristic. Taillard tested various types of neighbourhoods and the neighbourhood resulting from changing the position of one job proved to be the best.

Ogbu and Smith (1990a) proposed a Simulated Annealing approach to the PFSP which involved an initialisation with the Palmer and Dannenbring heuristics. The choice of a large neighbourhood and an acceptance probability function independent of the change in the makespan of the schedule resulted in near-optimal schedules for the problems tested. In a later work (Ogbu and Smith, 1990b) compared this Simulated Annealing with that of Osman and Potts. The results show a tie between the two metaheuristics with a slight advantage on the part of Osman and Potts's one.

Werner (1993) constructed a fast iterative method which obtained very good results by generating restricted numbers of paths in a search neighbourhood built upon the path structure of a feasible schedule. It is interesting to remark that while local search or fast iterative methods operate in a very small neighbourhood, Werner's path algorithm operates with a rather large neighbourhood but only evaluates the most interesting solutions in each step.

Reeves (1993) modified the SPIRIT algorithm and introduced several enhancements, such as the initialisation by the NEH heuristic and the insertion neighbourhood. The author compared this algorithm with the SA of Osman and Potts showing better results.

Chen et al. (1995) developed a simple Genetic Algorithm for the PFSP with various enhancements. The initial population is generated with the CDS and RA heuristics and also from simple job exchanges of some of the individuals. Only the crossover operator is applied (no mutation), and the crossover used is the *Partially mapped crossover* or PMX. Reeves (1995) also developed a Genetic Algorithm. In this case, the offspring generated in each step of the algorithm do not replace their parents but individuals from the generation that have a fitness value below average. Reeves used a crossover operator called C1, which is equivalent to the *One Point Order Crossover*. Another remarkable feature of the algorithm is that it uses an adaptive mutation rate. The algorithm uses a *shift* mutation which simply changes the position of one job. Reeves also chose to seed the initial population with a good sequence among randomly generated ones. This good sequence is obtained with the NEH heuristic. Also, the selection of the parents is somewhat different from what is common in genetic algorithms; parent 1 is selected using a fitness rank distribution whereas parent 2 is chosen using a uniform distribution. This Genetic Algorithm was one of the first methods to be tested against Taillard (1993)'s famous benchmark instances.

Ishibuchi et al. (1995) presented two Simulated Annealing algorithms characterized by having robust performance with respect to the temperature cooling schedule. Their results showed that their algorithms were comparable to the SA of Osman and Potts. Zegordi et al. (1995) demonstrated a hybrid technique by introducing problem domain knowledge into a Simulated Annealing algorithm. Their algorithm, called SA-MDJ, uses a "*Move desirability for*

Jobs” table which incorporates several rules that facilitate the annealing process. This way, fewer control parameters are needed. The algorithm is compared with Osman and Potts’s SA algorithm proving to be slightly inferior but much faster in the instances tested.

The Tabu Search of Moccellini (1995) is mainly based on Widmer and Hertz’s SPIRIT heuristic. The only difference exists in the step of calculating the initial solution for the problem. The analogy with the Travelling Salesman Problem is maintained, but the way the distance between jobs is calculated differs from that of Widmer and Hertz. Also, how the TSP problem is solved is different, since Moccellini uses the farthest Insertion Travelling Salesman Procedure (FITSP). The rest of the algorithm is essentially similar to SPIRIT.

The genetic algorithms proposed by Murata et al. (1996) use the two-point crossover operator and a shift mutation along with an elitist strategy to obtain good solutions for the PFSP. The algorithm performed worse than implementations of Tabu Search, Simulated Annealing and local search, so the authors implemented two hybrid versions of the Genetic Algorithm; Genetic Simulated Annealing and Genetic Local Search. In these algorithms, an “improvement” phase is performed before selection and crossover and this improvement is made with local search and simulated annealing algorithms respectively. The hybrid algorithms performed better than the non hybrid Genetic Algorithm, implementations of Tabu Search, Simulated Annealing and local search.

Nowicki and Smutnicki (1996) proposed a Tabu Search metaheuristic where once again, a reduced part of the neighbourhood is evaluated along with a fast method for obtaining the makespan. The neighbourhood is reduced with the idea of *blocks of jobs*, where jobs are clustered and the movements are made on a block-by-block basis instead of just moving single jobs. This

algorithm has proved to be one of the finest metaheuristics for the PFSP and the authors were then able to lower the best upper bounds for Taillard's benchmark instances.

Another hybrid Genetic Algorithm is that of Reeves and Yamada (1998), the idea behind this algorithm is the *Multi-step crossover fusion* or MSXF operator which coalesces a crossover operator with local search. The MSXF operator carries a biased local search from one parent using the other parent as a reference. The MSXF operator is defined by a neighbourhood structure and a measure of distance, which has some similarities with simulated annealing but with a constant temperature parameter. The results given outperform other metaheuristics and the authors also found new lower upper bounds for Taillard's benchmark instances. Another metaheuristic is given by Stützle (1998), and is given the name "Iterated Local Search" or ILS. In some way or another, the ideas behind ILS are those of descent search, local search, Simulated Annealing or Tabu Search; An initial solution is obtained and a *LocalSearch* procedure is executed over this initial solution, then a *Modify* procedure that slightly modifies the current solution is carried out. Then, *LocalSearch* is called again and the resulting schedule is deemed as the new current solution only if a certain acceptance criterion is met (procedure *AcceptanceCriterion*). The ILS procedure may also work with a list called *history* that works similarly to the tabu list in Tabu Search. The experimental results given in the paper show some very interesting conclusions. According to the tests conducted by Stützle, the ILS algorithm is much better than the Tabu Search of Taillard (1990) and also better than the Tabu Search of Nowicki and Smutnicki (1996).

Ben-Daya and Al-Fawzan (1998) implemented a Tabu Search algorithm with some extra features such as intensification and diversification schemes that provide better moves in the tabu search process. The algorithm proposed

provides similar results as the TS of Taillard and slightly better results than the SA of Ogbu and Smith. Moccellini and dos Santos (2000) presented a hybrid Tabu Search-Simulated Annealing heuristic that is compared with simple Tabu Search and simple Simulated Annealing implementations from the same authors, showing advantages for the hybrid approach.

Ponnambalam et al. (2001) evaluate a Genetic Algorithm which uses the GPX crossover (Generalised Position Crossover) and other features like shift mutation and a randomized initial solution.

Wodecki and Bożejko (2002) have recently proposed a SA algorithm that is run in a parallel computing environment. The authors compare the proposed algorithm with the NEH heuristic, the former showing better results. Finally, Wang and Zheng (2003) have presented a complex hybrid genetic algorithm in which the initialisation of the population uses the NEH heuristic and the mutation operator is replaced by a SA algorithm. The proposed algorithm also uses multicrossover operators that are applied to subpopulations in the original population.

4 Comparative Evaluation of Heuristics and Metaheuristics

Table 1 shows a summary of the different constructive and improvement heuristics reviewed in chronological order, and Table 2 shows information about the reviewed metaheuristics. In order to evaluate the different heuristics, a broad number of them have been coded, these are: Johnson’s algorithm (extended to the m machine case) (Johns), Page’s pairing sorting algorithm (Page), Palmer’s “slope index” heuristic (Palme), the Campbell et al. and Gupta’s heuristics (CDS and Gupta). The three heuristics by Dannenbring (RA, RACS and RAES), Nawaz et al. NEH heuristic (NEH), Hundal and

Year	Author/s	Acronym	Type*	Comments
1954	Johnson	Johns	C	exact for two machine case
1961	Page	Page	C	based on sorting
1964	Dudek and Teuton		C	based on Johnson's rule
1965	Palmer	Palme	C	based on slope indexes
1970	Campbell et al.	CDS	C	based on Johnson's rule
1971	Gupta	Gupta	C	based on slope indexes
1972	Gupta		C	three heuristics considered
1976	Bonney and Gundry		C	based on slope matching
1977	Dannenbring	RA,RACS, RAES	C/ I	three heuristics considered: RA, RACS, and RAES
1980	King and Spachis		C	5 dispatching rule based heuristics
1982	Stinson and Smith		C	6 heuristics, based on TSP
1983	Nawaz et al.	NEH	C	job priority/insertion
1988	Hundal and Rajgopal	HunRa	C	Palmer's based
1991	Ho and Chang	HoCha	I	between-jobs gap minimisation
1993	Sarin and Lefoka		C	last machine idle time minimisation
1998	Koulamas	Koula	C/ I	two phases, 1st Johnson based, 2nd phase improvement by job passing
2000	Suliman	Sulim	I	job pair exchange
2001	Davoud Pour	Pour	C	job exchanging
2003	Framinan et al.		C	study on the NEH heuristic

*C: Constructive, I: Improvement

Table 1: Constructive and improvement heuristics for the Permutation Flow-shop Problem.

Rajgopal's extension to Palmer's heuristic (HunRa), Ho and Chang's gap improvement heuristic (HoCha), Koulamas' two-phase heuristic (Koula), Suliman's heuristic (Sulim) and finally Davoud Pour's algorithm (Pour).

The metaheuristics implemented are: Osman and Potts' SA algorithm (SAOP), Widmer and Hertz's SPIRIT procedure (Spirit), Chen et al.'s GA algorithm (GAChen), Reeves' (GAReev) GA algorithm, also the hybrid GA + Local Search by Murata et al. (GAMIT), Stützle's Iterated Local Search (ILS) and finally the GA by Ponnambalam et al. (GAPAC). As we can see, the methods selected for the comparison are the most known heuristics and metaheuristics and also some original methods that are either recent and have not been

Year	Author/s	Acronym	Type	Comments
1989	Osman and Potts	SAOP	SA	
	Widmer and Hertz	Spirit	TS	initial solution based on the OTSP
1990	Taillard		TS	
	Ogbu and Smith		SA	
1993	Werner		other	path algorithms
	Reeves		TS	
1995	Chen et al.	GAChen	GA	PMX crossover
	Reeves	GAReev	GA	adaptive mutation rate
	Ishibuchi et al.		SA	two SA considered
	Zegordi et al.		SA	combines sequence knowledge
	Moccellin		TS	based on SPIRIT
1996	Murata et al.	GAMIT	hybrid	GA+Local Search/SA
	Nowicki and Smutnicki		TS	neighbourhood by blocks of jobs
1998	Stützle	ILS	other	Iterated Local Search
	Ben-Daya and Al-Fawzan		TS	intensification + diversification
	Reeves and Yamada		GA	GA operators with problem knowledge
2000	Moccellin and dos Santos		hybrid	TS + SA
2001	Ponnambalam et al.	GAPAC	GA	GPX crossover
	Wodecki and Bożejko		SA	parallel simulated annealing
2003	Wang and Zheng		hybrid	GA + SA, multicrossover operators

Table 2: Metaheuristics for the Permutation Flowshop Problem.

evaluated before or some that incorporate new ideas not previously used by other algorithms.

All efforts to code the complex TSAB algorithm of Nowicki and Smutnicki (1996) were unsuccessful. We tried several times to contact the authors but were unable to obtain the code for help. However, Stützle showed his ILS algorithm to be superior to the published results of the TSAB. Note that, however, the implementations that have been carried out in this work come from the explanations given in the original papers which could sometimes differ slightly from the original authors' implementations. All algorithms and methods (25 total) are coded in Delphi 6.0 and run in an Athlon XP 1600+ computer (1400 MHz) with 512 MBytes of main memory.

In order to make a fair comparison, the stopping criterion for all the meta-heuristics tested is a maximum of 50,000 makespan evaluations where every algorithm is run five independent times to obtain a final average of the results. For comparison reasons, we have also included a random procedure that evaluates random schedules and returns the best one (RAND). Furthermore, three common dispatching rules have been implemented, these are FCFS (First Come, First Served), SPT (Shortest Processing Time First) and LPT (Longest Processing Time First). The reason behind this is that dispatching rules are commonly used in practice. Note that SPT and LPT rules can produce non-permutation schedules where job passing is permitted.

4.1 Evaluation using Taillard’s instances

The instances of Taillard (1993) form a set of 120 problems of various sizes, having 20, 50, 100, 200 and 500 jobs and 5, 10 or 20 machines. There are 10 problems inside every size set. In total there are 12 sets and these are: 20x5, 20x10, 20x20, 50x5, 50x10, 50x20, 100x5, 100x10, 100x20, 200x10, 200x20 and 500x20. These problems are extremely difficult to solve and are a good benchmark to test the different methods reviewed. As of today, there are still 33 instances whose optimal solution has not been established.

For Taillard (1993)’s 120 problems we are going to work with the following performance measure:

$$\% \text{ Increase Over Optimum} = \frac{Heu_{sol} - Opt_{sol}}{Opt_{sol}} \cdot 100$$

Where Heu_{sol} is the solution obtained by a given algorithm and Opt_{sol} is the given optimum solution for each instance in the OR Library (<http://mscmga.ms.ic.ac.uk/jeb/orlib/flowshopinfo.html>) or the lowest known upper bound if the optimum for that instance is still unknown. The results

of the different heuristics can be seen in Table 3. Note that for every problem size we have averaged the 10 corresponding instances. As is shown in Ta-

Problem	RAND	FCFS	SPT	LPT	Johns	Page	Palme	CDS	Gupta
20x5	5.79	24.98	20.44	47.66	12.78	15.15	10.58	9.54	12.45
20x10	9.47	28.77	24.74	46.38	19.97	20.43	15.28	12.13	24.48
20x20	7.76	21.43	23.17	34.27	16.47	16.18	16.34	9.64	22.53
50x5	4.74	15.32	16.21	48.33	10.43	10.14	5.34	6.10	12.43
50x10	14.00	25.05	26.51	51.23	21.90	20.47	14.01	12.98	22.05
50x20	16.31	29.59	29.38	48.08	22.81	23.12	15.99	13.85	23.66
100x5	3.70	13.63	17.98	45.19	6.84	7.98	2.38	5.01	6.02
100x10	10.31	20.92	24.63	47.27	15.01	15.79	9.20	9.15	15.12
100x20	16.28	25.36	28.93	47.01	21.15	21.68	14.41	13.12	22.68
200x10	8.39	15.67	20.25	45.54	11.47	12.74	5.13	7.38	11.80
200x20	15.32	22.10	25.94	51.10	18.93	19.43	13.17	12.08	19.67
500x20	11.45	15.99	23.04	48.01	14.07	14.05	7.09	8.55	13.66
Average	10.29	21.57	23.44	46.67	15.99	16.43	10.74	9.96	17.21

Problem	RA	RACS	RAES	NEH	HunRa	HoCha	Koula	Sulim	Pour
20x5	8.86	7.71	4.95	3.35	9.35	6.94	7.68	4.46	12.05
20x10	15.40	10.66	8.62	5.02	13.34	10.51	11.82	7.84	12.34
20x20	16.35	8.16	6.41	3.73	13.47	8.30	11.89	6.69	10.71
50x5	6.30	5.40	3.28	0.84	4.21	3.33	4.03	2.20	6.58
50x10	15.05	12.19	10.41	5.12	13.35	11.29	12.13	8.46	14.44
50x20	18.88	12.86	10.00	6.20	14.99	12.40	14.93	9.62	14.87
100x5	3.54	4.58	3.25	0.46	1.99	2.70	3.12	1.28	4.06
100x10	10.48	8.72	7.31	2.13	8.12	7.96	7.50	5.75	7.82
100x20	16.96	12.48	10.56	5.11	13.65	11.10	14.04	9.28	13.18
200x10	6.17	7.11	6.16	1.43	4.50	5.11	5.09	4.09	5.52
200x20	12.67	11.83	10.39	4.37	12.59	9.99	11.60	8.85	11.50
500x20	8.34	8.38	7.77	2.24	6.75	7.14	6.82	6.06	7.69
Average	11.58	9.17	7.43	3.33	9.69	8.06	9.22	6.21	10.07

Table 3: Average percentage increase over the best solution known for the heuristic algorithms.

ble 3, the worst performing algorithms are the three dispatching rules, with the LPT being close to a 50% increase over the best known solution. The other two dispatching rules achieve over 21% (FCFS) and over 23% (SPT) increase respectively. This is an interesting result, since dispatching rules are the most used algorithms in professional sequencing and scheduling software, despite their very poor performance. We would like to point out that it is even better to use the RAND algorithm than to use dispatching rules by

a significant margin, so there is no reason for professional software to keep using these kinds of methods.

For a better analysis of the results we have performed a design of experiments and an analysis of variance (ANOVA) (see Montgomery, 2000) where we consider the different algorithms as a factor and all the instances solved as treatments. We have chosen to take the three dispatching rules out of the experiment since it is clear that their performance is very poor. This way we have 15 different algorithms and 1,800 treatments. It is important to note that in an ANOVA it is necessary to check the model's adequacy to the data and there are three important hypotheses that should be carefully checked: normality, homogeneity of variance (or homocedasticity) and independence of the residuals. We carefully checked all these three hypotheses and we found no basis for questioning the validity of the experiment.

The results indicate that there are statistically significant differences between the different algorithms with a p-value very close to zero. Figure 1 shows the means plot with LSD intervals (at the 95% confidence level) for the different algorithms. After analysing the means and intervals we can form nine homogeneous groups where within each group no statistically significant differences can be found. For example, the two lowest performers, the heuristics of Gupta and Page are, at the 95% confidence level, similar. The same applies for the Johnson rule and Page's heuristic. After these three heuristics we can see that the RA method and Palmer's heuristic are similar, and at the same time, this last algorithm is similar to many other methods (RAND, Pour, CDS and HunRa).

Interestingly enough, Gupta's heuristic is over 17% increase, despite the author's claims about being superior to Palmer's index algorithm (the author used a different benchmark from us). Moreover, Palmer's heuristic is superior to Gupta's in every problem size and it is clear from the experiment that

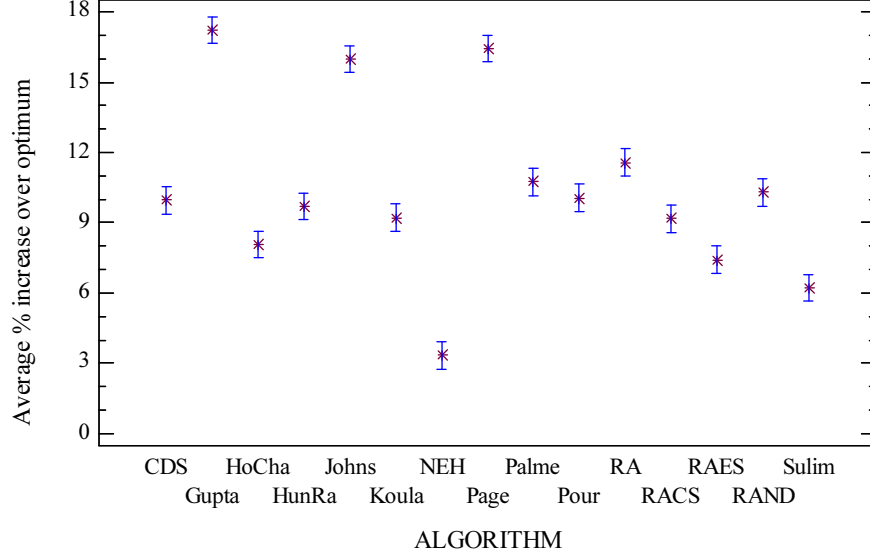


Figure 1: Means plot and LSD intervals for the Algorithm factor in the experiment.

there are statistically significant differences between the two methods.

In a better position are the RACS algorithm and the improvement heuristic of Ho and Chang (HoCha), although there is no clear difference between them. The RAES algorithm is better than the RACS with a 7.42% average percentage increase over the best solution.

The best two heuristics are the improvement method of Suliman and the NEH heuristic, although there are differences between the two which favour the second. Clearly, the NEH heuristic is, by far, the best heuristic among those tested. In some instances it even manages a mere increase of under 1% which is significantly better than every other method.

Some interesting results can be drawn from the study, first, among the most recent heuristics, Koulamas's two phase method has an average performance, despite being a remarkably sophisticated heuristic and permitting job-passing, thus allowing potentially better schedules. From the study no

statistically significant differences were obtained between this method and other older (and simpler) algorithms like the CDS or the RAND rule.

Davoud Pour’s (Pour) heuristic performance can also be compared to the CDS or to the Hundal and Rajgopal (HunRa) heuristics. This is a negative result since the Hundal and Rajgopal heuristic evaluates only three different schedules and Davoud Pour’s evaluates a total of $[n(n+1)/2] - 1$ schedules which is very slow for big instances. The results obtained by the author claiming to be a better heuristic than NEH, clearly do not hold with this benchmark.

In order to analyse the influence of the number of jobs (n) and number of machines (m) in the effectiveness of all algorithms we conducted a second experimental design where n and m are also considered as factors. As with the algorithm factor, n and m turned out to be statistically significant with real low p-values. In this case, in order to study the different interactions, the instances of more than 100 jobs had to be taken out of the experiment to maintain the necessary orthogonality in the experiment (i.e. there are no 200x5, 500x5 500x10 instances in Taillard’s benchmark). In Figure 2 we can see the effect that the number of machines has over all algorithms. As we can see, more machines lead to worse results, meaning that generally, instances with 20 machines are harder than instances with only five or 10 machines. No clear trend was observed for the factor n . There are, however, some interesting interactions, such as the one concerning the different algorithms and n that is shown in Figure 3. It can be observed that the NEH heuristic performs very well regardless of the number of jobs involved. Some algorithms, such as HoCha, CDS or Sulim are robust and show a statistically similar performance no matter how many jobs, whereas Johnson’s rule or Palmer’s heuristic are sensitive to the number of jobs. The situation is very similar for the interaction between the number of machines (m) and the different

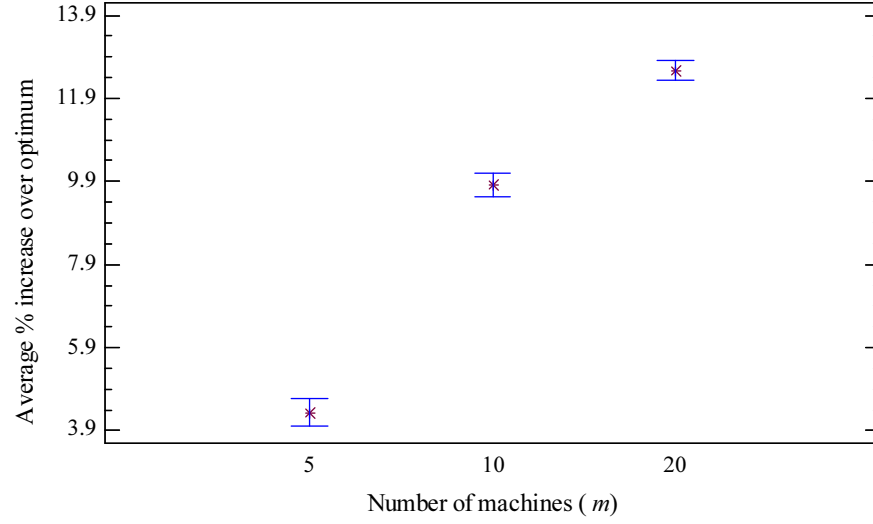


Figure 2: Means plot and LSD intervals for the number of machines factor in the experiment.

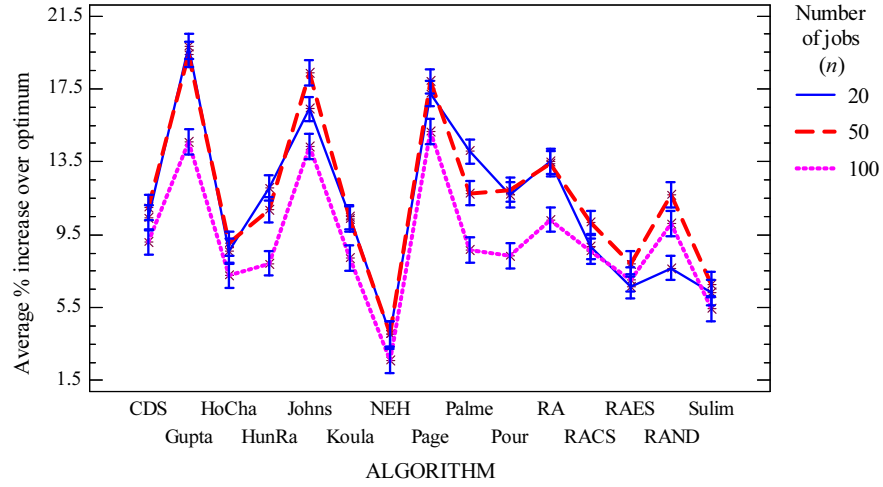


Figure 3: Means plot and LSD intervals for the interaction between the algorithms and the number of jobs (n).

algorithms, where the best results for the different methods are obtained in the five machine instances. Finally, it is worth noting the interaction between the factors (n) and (m), which is shown in Figure 4. It is clear that the most




Figure 4: Means plot and LSD intervals for the interaction between the number of jobs (n) and the number of machines (m).

difficult instances are those with 50 jobs and 20 machines and those with 100 jobs and 20 machines. As a matter of fact, these 20 instances are still open and no optimal solution has been established so far.

It is also interesting to analyse the efficiency of the different heuristics when evaluating the different instances. We have measured the time (in seconds) that a given method needs in order to provide a solution. The results for the classical heuristics are straightforward, the RAND, FCFS, SPT, LPT, Johnson, Page, Palmer, CDS, Gupta, RA and RACS heuristics all needed less than 0.5 seconds on average for all problem sizes. For the remaining heuristics, the times were similar for problem sizes up to 100x5. From that point, the times needed for solving the instances started to increase (Table 4). We can see that the RAES method is very fast, needing less than five seconds on average but only for the biggest instances. The NEH and Hundal and

Problem	RAND	RAES	NEH	HunRa	HoCha	Koula	Sulim	Pour
100x5	0.58	<0.5	<0.5	<0.5	<0.5	<0.5	<0.5	0.86
100x10	1.08	<0.5	<0.5	<0.5	<0.5	<0.5	<0.5	1.54
100x20	2.02	<0.5	<0.5	<0.5	<0.5	0.70	<0.5	2.92
200x10	2.26	<0.5	<0.5	<0.5	<0.5	1.04	<0.5	13.81
200x20	4.50	<0.5	<0.5	<0.5	<0.5	4.52	<0.5	26.00
500x20	22.38	4.67	<0.5	<0.5	1.18	64.55	13.61	532.61
Average	2.95	<0.5	<0.5	<0.5	<0.5	5.93	1.20	48.20

Table 4: Average CPU time in seconds used by the different heuristics.

Rajgopal, as well as Ho and Chang are also very efficient. It is important to note that the NEH heuristic has been implemented with the speed improvement suggested by Taillard (1990), otherwise, the classical NEH without this improvement, is very slow and almost impracticable for large instances. The most recent constructive and improvement heuristics, such as Koulamas’s method take more than one minute in the 500x20 instances, whereas Suliman’s is faster with less than 20 seconds.

Davoud Pour’s heuristic times grow exponentially with the size of the problem considered, reaching almost nine minutes on average for the biggest instances. This is due to the extensive work needed in each iteration. We think these time requirements are unbearable, especially considering this method’s low performance. In this situation it is much better to use the simple RAND method, which needs little more than 20 seconds to evaluate 50,000 random makespans and the solution obtained is of the same quality as that of Davoud Pour’s heuristic.

Now we are going to analyse the effectiveness and efficiency of the different metaheuristics reviewed. The results for the metaheuristics can be seen in Table 5. The first conclusion that could be drawn is that the ILS and GAReev methods are the most effective and the GA of Ponnambalam et al. (GAPAC) the least. However, the CPU times range from little more than four seconds on average for the Spirit method to more than half a minute for

Problem	SAOP	Spirit	GAChen	GAReev	GAMIT	ILS	GAPAC
20x5	1.39 (<0.5)	5.22 (<0.5)	3.82 (<0.5)	0.70 (1.08)	4.21 (<0.5)	0.24 (4.01)	8.98 (<0.5)
20x10	2.66 (<0.5)	5.86 (<0.5)	4.89 (<0.5)	1.92 (1.17)	5.40 (<0.5)	0.77 (4.09)	13.61 (<0.5)
20x20	2.31 (<0.5)	4.58 (<0.5)	4.17 (0.60)	1.53 (1.35)	4.53 (<0.5)	0.85 (4.63)	11.03 (0.57)
50x5	0.69 (<0.5)	2.03 (<0.5)	2.09 (0.77)	0.26 (1.82)	3.11 (<0.5)	0.12 (6.38)	6.50 (0.58)
50x10	4.25 (0.60)	5.88 (0.52)	6.60 (1.00)	2.58 (2.08)	8.38 (0.52)	2.01 (9.94)	16.41 (0.83)
50x20	5.13 (1.04)	7.21 (0.97)	8.03 (1.45)	3.76 (2.53)	10.65 (0.96)	3.29 (11.82)	18.56 (1.28)
100x5	0.40 (0.60)	1.06 (0.53)	1.32 (1.79)	0.18 (3.97)	5.41 (0.52)	0.11 (15.31)	5.32 (1.07)
100x10	1.88 (1.10)	5.07 (1.03)	3.75 (2.26)	1.08 (4.49)	12.05 (1.02)	0.66 (18.79)	12.34 (1.57)
100x20	5.21 (2.09)	10.15 (2.00)	7.94 (3.24)	3.94 (5.54)	18.24 (1.99)	3.17 (24.04)	18.25 (2.55)
200x10	1.56 (2.29)	9.03 (2.25)	2.70 (5.97)	0.82 (12.91)	7.52 (2.20)	0.49 (33.73)	9.75 (3.24)
200x20	4.83 (4.59)	16.17 (4.51)	7.07 (8.18)	3.33 (15.16)	15.35 (4.50)	2.74 (41.80)	17.06 (5.46)
500x20	3.40 (39.48)	13.57 (39.70)	4.61 (55.30)	1.83 (101.71)	12.17 (37.82)	1.29 (192.03)	12.61 (30.62)
Average	2.81 (4.42)	7.15 (4.38)	4.75 (6.77)	1.83 (12.82)	8.92 (4.21)	1.31 (30.55)	12.53 (4.04)

Table 5: Average percentage increase over the best solution known and average CPU time in seconds (between parenthesis) for the metaheuristic algorithms.

the ILS algorithm.

Considering the number of makespan evaluations as the termination criterion along with the elapsed times, facilitates the extrapolation of results among different computing platforms. But since all algorithms have been coded in the same programming language and are run on the same computer, we can establish a better comparison by fixing a given amount of time as a termination criterion. We have tested all metaheuristics by considering the following expression: $\text{termination criterion} = (n \cdot 60,000) / (1,000)$ milliseconds. This is the same as saying that we allow 60 seconds for every 1,000 jobs in the problem. Thus we will allow a maximum time of 1,200 milliseconds for the 20 job instances and 30,000 milliseconds for the biggest 500 job instances. From the Table 5 we see that the factor that affects CPU time the most is the number of jobs so increasing the time along with the number of jobs seems natural. The results (averaged for five independent runs) are shown in Table 6. We see that now the differences between the methods ILS and GAReev have narrowed. Again, to compare the results we have designed an experiment considering the metaheuristic algorithms as a factor. Note that the

Problem	SAOP	Spirit	GAChen	GAReev	GAMIT	ILS	GAPAC
20x5	1.47	2.14	3.67	0.71	3.28	0.29	8.26
20x10	2.57	3.09	5.03	1.97	5.53	1.26	12.88
20x20	2.22	2.99	4.02	1.48	4.33	1.04	10.62
50x5	0.52	0.57	2.31	0.23	1.96	0.12	7.19
50x10	3.65	4.08	6.65	2.47	6.25	2.38	16.53
50x20	4.97	4.78	7.92	3.89	7.53	4.19	18.84
100x5	0.42	0.32	1.18	0.18	1.33	0.12	5.21
100x10	1.73	1.69	3.91	1.06	3.66	0.85	12.38
100x20	4.90	4.49	7.82	3.84	9.70	3.92	17.88
200x10	1.33	0.95	2.70	0.85	6.47	0.54	9.94
200x20	4.40	3.85	7.01	3.47	14.56	3.34	16.83
500x20	3.48	2.17	5.62	1.98	12.47	1.82	12.67
Average	2.64	2.60	4.82	1.84	6.42	1.65	12.44

Table 6: Average percentage increase over the best solution known for the metaheuristic algorithms. Maximum elapsed time stopping criterion.

GAPAC algorithm gives very poor results and has been disregarded in the experiment. Again, the ANOVA results indicate that there are statistically significant differences between the algorithms considered, with the p-value being almost zero. To evaluate which algorithms perform better than others we focus on the means plot shown in Figure 5. We can see that there are four clear homogeneous groups, (ILS, GAReev), (SAOP, Spirit), (GAChen) and (GAMIT). The two best metaheuristics are the ILS and the GAReev methods and the worst is the GAMIT algorithm. We consider the result of the GAMIT algorithm rather poor as it is a complex hybrid algorithm. Also, it is worthwhile to note that the GAPAC’s effectiveness is even lower than the much more simplistic RAND rule. This conclusion is different to that of the authors’, who claimed their algorithm to be better than the NEH or the CDS heuristics. This is clearly not the case for the benchmarks and evaluations shown here. SAOP and Spirit’s effectiveness is very similar although the implementation of the SAOP method is much simpler. Finally, we think that the ILS method compares better since, even though it is similar to the GAReev in terms of effectiveness, it is very simple to code, and it is much simpler than a Genetic Algorithm. Furthermore, the ILS algorithm can be

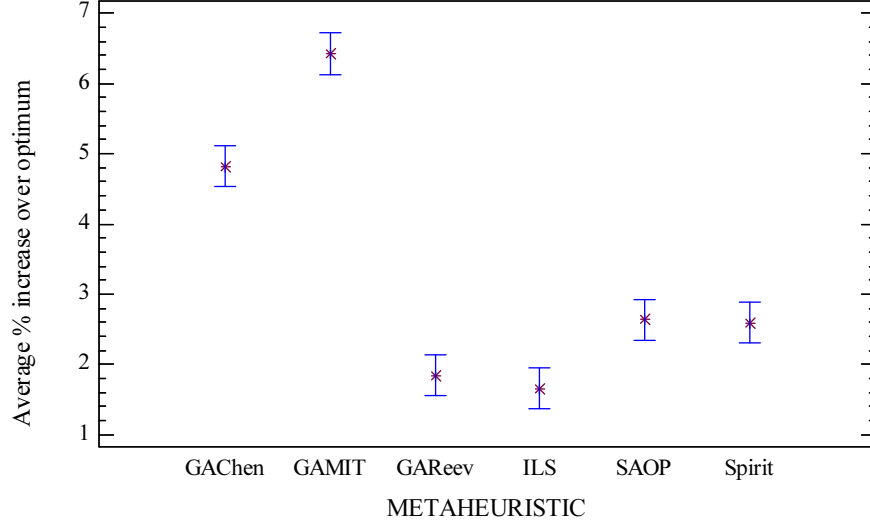


Figure 5: Means plot and LSD intervals for the Metaheuristic factor in the experiment.

coupled with a more powerful Local Search phase than the one used here. If we analyse again the influence of the number of jobs (n) and number of machines (m) in the effectiveness of the different metaheuristics we find almost identical results to those obtained when evaluating the heuristic methods. For example, in Figure 6 we can observe the interaction between the different metaheuristics and the number of jobs. As we can see, the profile of the different algorithms is almost the same regardless of the number of jobs in the instance, although there are some aspects worth noting. For example, the SAOP algorithm is more adequate than the Spirit for instances with 20 jobs and the ILS algorithm is also better than the GAReev for the instances with this same number of jobs.

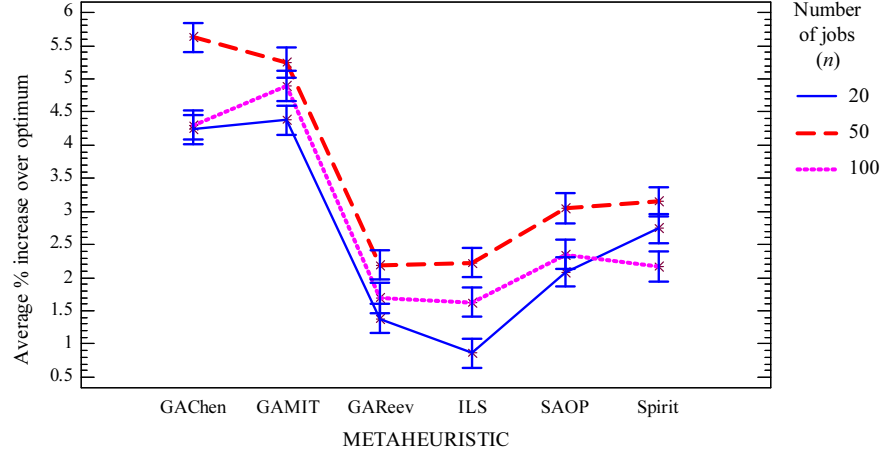


Figure 6: Means plot and LSD intervals for the interaction between the metaheuristics and the number of jobs (n).

5 Conclusions

In this paper we have given an extensive review and evaluation of many heuristics and metaheuristics for the permutation flow shop scheduling problem, also known as $n/m/P/F_{max}$ or as $F/prmu/C_{max}$. A total of 25 algorithms have been coded and tested against Taillard's (1993) famous 120 instance benchmark.

As a conclusion, we can now say that Nawaz et al.'s (1983) algorithm (NEH) is the best heuristic for Taillard's benchmarks. The performance of this heuristic, if implemented as in Taillard (1990) is outstanding, with CPU times below 0.5 seconds even for the largest 500x20 instances.

From the metaheuristics tested, Stützle's (1998) Iterated Local Search (ILS) method, and the genetic algorithm of Reeves are better than the other algorithms evaluated.

We would like to point out that dispatching rules have shown a very poor performance in Taillard's benchmarks, worse even than a simplistic RAN-

DOM rule, by a significant margin. However, these rules are still used in practice in professional software and by production managers in many industrial sectors. More elaborated heuristics outperform dispatching rules to a significant degree.

We have also shown that Tabu Search, Simulated Annealing and Genetic Algorithms are good metaheuristics for this problem, although the latter type needs a good initialisation of the population and/or problem domain knowledge in order to achieve a good performance. This result leads to the same conclusion reached by Murata et al. (1996).

Acknowledgments

This work is funded by the Polytechnic University of Valencia, Spain, under an interdisciplinary project, and by the Spanish Department of Science and Technology (research project ref. DPI2001-2715-C02-01).

The authors would like to thank the two anonymous referees for their insightful comments in a previous version of this paper and also to the ACLE office.

References

- Ashour, S. (1970). An Experimental Investigation and Comparative Evaluation of Flow-Shop Scheduling Techniques. *Operations Research*, 18(3):541–549.
- Baker, K. R. (1974). *Introduction to Sequencing and Scheduling*. John Wiley & Sons, New York.
- Ben-Daya, M. and Al-Fawzan, M. (1998). A tabu search approach for the flow shop scheduling problem. *European Journal of Operational Research*, 109:88–95.
- Bonney, M. and Gundry, S. (1976). Solutions to the Constrained Flowshop Sequencing Problem. *Operational Research Quarterly*, 27(4):869–883.

- Byung Park, Y., Pegden, C. D., and Ensore, E. E. (1984). A survey and evaluation of static flowshop scheduling heuristics. *International Journal of Production Research*, 22(1):127–141.
- Campbell, H. G., Dudek, R. A., and Smith, M. L. (1970). A Heuristic Algorithm for the n Job, m Machine Sequencing Problem. *Management Science*, 16(10):B630–B637.
- Chen, C.-L., Vempati, V. S., and Aljaber, N. (1995). An application of genetic algorithms for flow shop problems. *European Journal of Operational Research*, 80:389–396.
- Conway, R. W., Maxwell, W. L., and Miller, L. W. (1967). *Theory of Scheduling*. Addison-Wesley, Reading, MA.
- Dannenbring, D. G. (1977). An Evaluation of Flow Shop Sequencing Heuristics. *Management Science*, 23(11):1174–1182.
- Davoud Pour, H. (2001). A new heuristic for the n -job, m -machine flow-shop problem. *Production Planning and Control*, 12(7):648–653.
- Dudek, R. A., Panwalkar, S. S., and Smith, M. L. (1992). The Lessons of Flowshop Scheduling Research. *Operations Research*, 40(1):7–13.
- Dudek, R. A. and Teuton, Jr, O. F. (1964). Development of m -Stage decision Rule for Scheduling n Jobs Through m Machines. *Operations Research*, 12(3):471–497.
- Framinan, J. M., Leisten, R., and Rajendran, C. (2003). Different initial sequences for the heuristic of Nawaz, Ensore and Ham to minimize makespan, idletime or flowtime in the static permutation flowshop sequencing problem. *International Journal of Production Research*, 41(1):121–148.
- Garey, M., Johnson, D., and Sethi, R. (1976). The Complexity of Flowshop and Jobshop Scheduling. *Mathematics of Operations Research*, 1(2):117–129.
- Graham, R., Lawler, E., Lenstra, J., and Rinnooy Kan, A. (1979). Optimization and Approximation in Deterministic Sequencing and Scheduling: A Survey. *Annals of Discrete Mathematics*, 5:287–326.
- Gupta, J. N. (1971). A Functional Heuristic Algorithm for the Flowshop Scheduling Problem. *Operational Research Quarterly*, 22(1):39–47.
- Gupta, J. N. D. (1972). Heuristic Algorithms for Multistage Flowshop

- Scheduling Problem. *AIIE Transactions*, 4(1):11–18.
- Ho, J. C. and Chang, Y.-L. (1991). A new heuristic for the n -job, M -machine flow-shop problem. *European Journal of Operational Research*, 52:194–202.
- Hundal, T. S. and Rajgopal, J. (1988). An Extension of Palmer’s Heuristic for the Flow Shop Scheduling Problem. *International Journal of Production Research*, 26(6):1119–1124.
- Ishibuchi, H., Misaki, S., and Tanaka, H. (1995). Modified simulated annealing algorithms for the flow shop sequencing problem. *European Journal of Operational Research*, 81:388–398.
- Johnson, S. (1954). Optimal Two- and Three-Stage Production Schedules with Setup Times Included. *Naval Research Logistics Quarterly*, 1:61.
- King, J. R. and Spachis, A. S. (1980). Heuristics for flow-shop scheduling. *International Journal of Production Research*, 18(3):345–357.
- Koulamas, C. (1998). A new constructive heuristic for the flowshop scheduling problem. *European Journal of Operational Research*, 105:66–71.
- Moccellin, J. a. V. (1995). A New Heuristic Method for the Permutation Flow Shop Scheduling Problem. *Journal of the Operational Research Society*, 46:883–886.
- Moccellin, J. a. V. and dos Santos, M. O. (2000). An adaptive hybrid meta-heuristic for permutation flowshop scheduling. *Control and Cybernetics*, 29(3):761–771.
- Montgomery, D. C. (2000). *Design and Analysis of Experiments*. John Wiley & Sons, fifth edition.
- Murata, T., Ishibuchi, H., and Tanaka, H. (1996). Genetic algorithms for flowshop scheduling problems. *Computers & Industrial Engineering*, 30(4):1061–1071.
- Nawaz, M., Enscoe, Jr, E. E., and Ham, I. (1983). A Heuristic Algorithm for the m -Machine, n -Job Flow-shop Sequencing Problem. *OMEGA, The International Journal of Management Science*, 11(1):91–95.
- Nowicki, E. and Smutnicki, C. (1996). A fast tabu search algorithm for the permutation flow-shop problem. *European Journal of Operational Research*, 91:160–175.
- Ogbu, F. and Smith, D. (1990a). The application of the simulated annealing

- algorithms to the solution of the $n/m/C_{max}$ flowshop problem. *Computers & Operations Research*, 17(3):243–253.
- Ogbu, F. and Smith, D. (1990b). Simulated Annealing for the Permutation Flowshop Problem. *OMEGA, The International Journal of Management Science*, 19(1):64–67.
- Osman, I. and Potts, C. (1989). Simulated Annealing for Permutation Flowshop Scheduling. *OMEGA, The international Journal of Management Science*, 17(6):551–557.
- Page, E. S. (1961). An Approach to the Scheduling of Jobs on Machines. *Journal of the Royal Statistical Society, B Series*, 23(2):484–492.
- Palmer, D. (1965). Sequencing Jobs through a Multi-Stage Process in the Minimum Total Time - A Quick Method of Obtaining a Near Optimum. *Operational Research Quarterly*, 16(1):101–107.
- Pinedo, M. (2002). *Scheduling: Theory, Algorithms and Systems*. Prentice Hall, New Jersey, second edition.
- Ponnambalam, S. G., Aravindan, P., and Chandrasekaran, S. (2001). Constructive and improvement flow shop scheduling heuristics: an extensive evaluation. *Production Planning and Control*, 12(4):335–344.
- Potts, C. N., Shmoys, D. B., and Williamson, D. P. (1991). Permutation vs. non-permutation flow shop schedules. *Operations Research Letters*, 10:281–284.
- Reeves, C. and Yamada, T. (1998). Genetic Algorithms, Path Relinking, and the Flowshop Sequencing Problem. *Evolutionary Computation*, 6(1):45–60.
- Reeves, C. R. (1993). Improving the Efficiency of Tabu Search for Machine Scheduling Problems. *Journal of the Operational Research Society*, 44(4):375–382.
- Reeves, C. R. (1995). A Genetic Algorithm for Flowshop Sequencing. *Computers & Operations Research*, 22(1):5–13.
- Reisman, A., Kumar, A., and Motwani, J. (1994). Flowshop Scheduling/Sequencing Research: A Statistical Review of the Literature, 1952–1994. *IEEE Transactions on Engineering Management*, 44(3):316–329.
- Sarin, S. and Lefoka, M. (1993). Scheduling Heuristic for the n -Job m -Machine Flow Shop. *OMEGA, The International Journal of Management*

- Science*, 21(2):229–234.
- Stinson, J. P. and Smith, A. W. (1982). A heuristic programming procedure for sequencing the static flowshop. *International Journal of Production Research*, 20(6):753–764.
- Stützle, T. (1998). Applying Iterated Local Search to the Permutation Flow Shop Problem. Technical report, AIDA-98-04, FG Intellektik, TU Darmstadt.
- Suliman, S. (2000). A two-phase heuristic approach to the permutation flowshop scheduling problem. *International Journal of production economics*, 64:143–152.
- Taillard, E. (1990). Some efficient heuristic methods for the flow shop sequencing problem. *European Journal of Operational Research*, 47:67–74.
- Taillard, E. (1993). Benchmarks for basic scheduling problems. *European Journal of Operational Research*, 64:278–285.
- Turner, S. and Booth, D. (1987). Comparison of Heuristics for Flow Shop Sequencing. *OMEGA, The International Journal of Management Science*, 15(1):75–78.
- Wang, L. and Zheng, D. Z. (2003). An Effective Hybrid Heuristic for Flow Shop Scheduling. *The International Journal of Advanced Manufacturing Technology*, 21:38–44.
- Werner, F. (1993). On the heuristic solution of the permutation flow shop problem by path algorithms. *Computers & Operations Research*, 20(7):707–722.
- Widmer, M. and Hertz, A. (1989). A new heuristic method for the flow shop sequencing problem. *European Journal of Operational Research*, 41:186–193.
- Wodecki, M. and Bożejko, W. (2002). Solving the Flow Shop Problem by Parallel Simulated Annealing. In Wyrzykowski, R., Dongarra, J., Paprzycki, M., and Waśniewski, J., editors, *Parallel Processing and Applied Mathematics, 4th International Conference, PPAM 2001*, volume 2328 of *Lecture Notes in Computer Science*, pages 236–244. Springer-Verlag.
- Zegordi, S. H., Itoh, K., and Enkawa, T. (1995). Minimizing makespan for flowshop scheduling by combining simulated annealing with sequencing knowledge. *European Journal of Operational Research*, 85:515–531.